

Cross-validated per-class MCC: baseline and SMOTE variants



Models: LR = logistic regression, M1 = 1-hidden MLP (64), M2 = 2-hidden MLP (128,64). Symbols: ‡ class-reweighted BCE, ★ focal loss, † ridge ($\lambda=1e-4$). All at MCC-optimized thresholds; ctrl = random-label control.